

Top Trading Cycles in a practical setting

Author: Lars Møen Haukland

Supervisor: Fredrik Manne



UNIVERSITETET I BERGEN
Fakultet for naturvitenskap og teknologi

June, 2026

Abstract

This thesis explores the Top Trading Cycles algorithm in the practical setting of the allocations of GPs, General Practitioner, in Norway for people wanting to switch or get a GP. It explores how different focus of *good* results can change both the runtime and results, and how to make the algorithm run as efficiently as possible. This thesis proposes an implementation of the Top Trading Cycles that can reduce the waitlists for GPs in Norway by over 80%.

Acknowledgements

I would like to thank my supervisor Fredrik Manne for his guidance and support throughout this project. I would also like to thank Juni Weisteen Bjerde for input and help during the project. In addition Juni and Mathias Bertntert together theorised the exact algorithms so that I could implement them.

Contents

1. Introduction	5
1.1. Thesis Outline	5
2. Background	7
2.1. Related problems	7
2.1.1. One-to-One assignment problems	7
2.1.1.1. Housing market	7
2.1.1.2. Kidney exchange	7
2.1.2. Many-to-One assignment problems	8
2.1.2.1. The College Admissions problem	8
2.2. Top Trading Cycles	8
2.3. Cycle cancelling	9
2.4. Related Work	9
3. Problem Formalization	10
3.1. The GP allocation problem	10
3.1.1. Input	10
3.1.2. Feasible solution	10
3.1.3. Optimization criterion	10
3.1.3.1. Lexicographic maximization by priority	11
3.1.3.2. Maximizing cardinality	11
3.1.3.3. Maximizing utility	12
3.1.4. Priority function	13
4. Implementation	15
4.1. Different graph representations used	15
4.1.1. Patient and Doctor graph	15
4.1.2. Doctor graph collapsed edges	15
4.2. Greedy DFS	15
4.3. Exact algorithm for maximizing total switches	17
4.3.1. Graph structure	17
4.3.2. Algorithm	17
4.4. Exact algorithm for maximizing priority	19
4.5. Metaheuristics	22
5. Experimental Results	23
5.1. Experimental Setup	23
5.2. Performance Comparison	23
5.3. Impact of Priority Strategies	23
6. Conclusion	24
6.1. Summary	24
6.2. Future Work	24
Bibliography	25
Index of Figures	26

1. Introduction

The Norwegian general practitioner (GP) system, started in 2001, has aim to serve the Norwegian population with a stable GP. A GP can follow his or hers patients over time, creating a trustful and safe relationship [1]. Today this system works as intended, and GPs handle most of the populations needs. As of December 2025 the number of GPs in Norway is 5720 and the total number of citizens in the Norwegian GP system is 5.6 million. The amount of patients each GP has varies but on average each GP has about a thousand patients [2].

One inefficiency in the GP system is when a person wants to switch from one GP to another. If the GP they want to switch to does not have any more capacity for new patients, the person has to sign up in a queue and wait for a free slot to open up with that GP. In December the number of open GP lists was 1340, indicating that 77% of all GPs had no extra capacity [3]. In a more densely populated city such as Bergen, the numbers are even worse. In December 2025 Bergen had 268 GPs and only 7 GPs that had available capacity for new patients. That's 97% of the GPs in Bergen had no extra capacity [3].

When a person wants to switch to a GP that has no excess capacity they must sign up and then wait in line. In December 2025 there were about 300,000 people waiting for GPs in Norway [2]. As there are only 5720 doctors each doctor has on average more than 50 people on his or hers waiting list. Since many are switching from a GP it follows that it is likely that there are cycles, in that a person P_A wants doctor D_A which currently has P_B as a patient, but P_B wants to switch to doctor D_B who currently has P_A as a patient. In this case patient P_A ends up waiting for an open slot with P_B 's doctor and vice versa, but logically instead of waiting P_A and P_B could just swap doctors. This can be generalized to longer "waiting cycles" that could all be serviced by allowing for simultaneous switching.

In this thesis, we study this problem and come up with solutions for how one can help people who are waiting in line so that they can get their preferred choice. In particular we propose that we can use cycle detecting algorithms, to find optimal switches between patients and massively reduce the number of people on waiting lists by using these algorithms periodically. Our algorithms are based on existing cycle cancelling and TTC algorithms, and focus on two measures of fairness:

- Switching as many patients as possible
- Switching as high priority patients as possible

We have developed and tested algorithms for these cases.

In addition we compare these optimal algorithms with the existing TTC algorithm and a greedy algorithm, comparing gain of good vs runtime.

1.1. Thesis Outline

The remainder of this thesis is organized as follows:

- **Section 2** provides background on the Top Trading Cycles mechanism, cycle cancelling algorithms and how these methods have been used in similar problems.
- **Section 3** formally defines the computational problems arising from TTC as variants of cycle cover problems on directed graphs.

- **Section 4** describes our Rust implementation and the algorithmic strategies employed.
- **Section 5** presents experimental results comparing different priority strategies.
- **Section 6** concludes and discusses future work.

2. Background

Problems like the GP allocation problem have been studied for quite some time with a lot of variations. In addition there are different perspectives on how to solve these problems. While in economics they focus on the fairness of an assignment, trying to make a stable and strategy proof assignment. In this chapter we describe similar problems for the GP allocation problem, describe the Top Trading Cycles algorithm (TTC) and Huitfeldt et.al implementation of TTC for the GP allocation problem. Finally we describe cycle cancelling, an optimization technique that our exact algorithms are based on.

2.1. Related problems

There are variations of assignment problems. Usually in assignment problems we mention agents and objects and agents want objects, in the GP allocation problem the patients would be agents and doctors would be objects. The typical assignment problems are One-to-One assignment problems, meaning one agent gets one object and an object can only be “owned” by one agent. In addition we also have Many-to-One assignment problems, like the GP allocation problem, one agent has one object but one object can be “owned” by more than one agent. It might be misleading to say that our patients “own” doctors but this makes more sense in other problems that we will describe, owning a doctor in our case means having that doctor as a current doctor.

2.1.1. One-to-One assignment problems

2.1.1.1. Housing market

The Housing Market model introduced in Shapley and Scarf [4] describes a model with traders (agents) and indivisible goods (objects). These goods can be houses, making it a housing market. Each agent already has one house but may want to switch. Every agent has a preference list ordering every house in order of preference, Shapley and Scarf combine all these preference lists to make a preference matrix [4].

It is in this article from Shapley and Scarf [4] that they also introduce the Top Trading Cycles algorithm and describe how it is fitting for models like the Housing Market. We also see similarities to the GP allocation problem in that we have agents already owning an object but wanting to switch. Two differences is the One-to-One since a doctor can have multiple patients compared to that one house can only be owned by one agent, and that in the Housing market each agent has a complete preference list while in the GP assignment problem each agent has only one preferred doctor.

2.1.1.2. Kidney exchange

The Kidney Exchange problem describes an issue when trying to find compatible kidney donors. It describes the case where we have pairs of patients and donors, but the patient is incompatible with the donors kidney. We then need to find some way to swap around the donors to compatible patients. Either with pairwise swapping or with larger cycles. If we visualize the problem as a directed graph, each vertex is a patient and donor pair. An edge from Pair A to Pair B means that Donor A is compatible with Patient B. Now cycles can form as in Figure 1 we have a cycle of length 3 $A \rightarrow B \rightarrow C$.




Figure 1: An example graph for the kidney exchange problem.

Abraham, Blum and Sandholm [5] show that, with unbounded cycle length, a maximum-cardinality and maximum-weight exchange can be found in polynomial time [6]. An issue often arising with unrestricted cycle length is if a donor suddenly refuses to donate, if this donor refuses after their patient has received a kidney then a patient is left without a donor and cannot exchange later. Because of this the problem is often considered with a restricted cycle length to allow all operations to be executed at the same time. This way no donor can refuse in the middle. For each pair in a cycle you need 2 operations, for a cycle of length k you need $2k$ operations at the same time. Finding a maximum-cardinality exchange with cycles of length at most k , for any fixed $k \geq 3$, is an NP-hard computational problem [5], [6].

The Kidney Exchange problem and the GP allocation problem have much in common. In both problems agents already have an object, but want to switch for another. We need to find cycles to exchange the objects in such a manner that most amount of people are happy. The key difference is that only one donor can be “owned” by one patient while in the GP allocation problem a doctor can be “owned” by multiple patients.

2.1.2. Many-to-One assignment problems

2.1.2.1. The College Admissions problem

In the College Admissions problem, introduced by Gale and Shapley, describes the problem of assigning applicants among colleges. Each college has a certain quote, a maximum number of applicants to admit and each applicant ranks the colleges in order of preference. An applicant can skip colleges he or she would never accept under any circumstances [7].

The problem is then to find a stable assignment meaning that there does not exist a case where there are two applicants α and β who are assigned to colleges A and B , respectively, although β prefers A to B and A prefers β to α . Gale and Shapley introduce an algorithm called deferred-acceptance (DA) to find the optimal stable assignment and prove that this runs in polynomial time and finds the optimal stable assignment [7].

There are close similarities between the College Admissions problem and the GP assignment problem. One difference is that patients, equivalent to the applicants, only prefer one doctor. Another difference is that patients already have existing doctors, making the problem different since applicants do not already have a college they want to switch from.

2.2. Top Trading Cycles

Top Trading Cycles (TTC), developed by Gale and published by Shapley and Scarf [4], is an algorithm to find a re-allocation of goods without using money. As mentioned it was introduced in an article together with the Housing Market problem. The algorithm finds a re-

allocation of houses to traders, such that all mutually-beneficial exchanges have been realized [8]. For the Housing Market problem the TTC algorithm does as follows [8]:

1. Ask each agent to indicate his “top” (most preferred) house.
2. Draw an arrow from each agent i to the agent, denoted $\text{Top}(i)$, who holds the top house of i .
3. Find a cycle (guaranteed to exist) and execute the trades in that cycle. Remove all involved agents from the graph.
4. If there are remaining agents, go back to step 1.

Note that a cycle is guaranteed to exist if there are agents still left. This is because of the pidgeonhole principle, since each agent has one outgoing edge we have to have a cycle. The cycle can also be of length 1 in which case an agent’s top house is its own, then we treat it as a normal cycle and remove the agent from the graph. For this classical TTC implementation to work we need to remember that agents have a strict preferences. Each agent has a complete list of all houses ranked in order of preference. Because of this as houses get re-allocated agents are bound to either find a trade cycle with others or end up having their $\text{Top}(i)$ house as their own.

Consider the example of 3 agents, A , B and C , each has a house and have a strict preference list. Lets go through how TTC would handle this:

The preference lists are

- $A = [B, C, A]$
- $B = [C, A, B]$
- $C = [B, C, A]$

If we make the graph where each agent point to their top we get Figure 2.




Figure 2: An example Housing Market for TTC.

We have the cycle $B \rightarrow C$ so we execute the trades, remove agents from the graph and create the edges again. The last remaining node is A and as B and C are not in the graph, $\text{Top}(i)$ of A is now A .




Figure 3: An example Housing Market for TTC, A ’s $\text{Top}(i)$ is itself.

We execute this trade and are left with no more agents making TTC terminate.

Roth proved that TTC is strategy proof, meaning all agents should prefer their true preferences [9].

2.3. Cycle cancelling

2.4. Related Work

3. Problem Formalization

In this chapter we define the GP allocation problem and also some variations of it. In addition we go through how to construct a graph given a GP allocation problem and finally how we can reduce the GP allocation problem to the Cycle Cover problem in directed graphs.

3.1. The GP allocation problem

3.1.1. Input

In the GP allocation problem we are given a set of patients $P = \{p_1, p_2, \dots, p_n\}$ and a set of doctors $D = \{d_1, d_2, \dots, d_m\}$. We are also given the current and preferred GP assignments, for each patient:

- $D_{\text{cur}}[i]$ denotes the index of the current doctor assigned to patient p_i (or $\perp$ if unassigned)
- $D_{\text{pref}}[i]$ denotes the index of the preferred doctor for patient p_i

In addition we are given a priority function $R : P \rightarrow \mathbb{N}$, mapping some positive integer to each patient, with higher numbers indicating a higher priority to help this patient.

3.1.2. Feasible solution

A feasible solution for the GP allocation problem is a subset of patients $S \subseteq P$ such that the directed graph

$$G_S = (S, E_S), E_S = \{(a, b) \mid a, b \in S, D_{\text{pref}}[\text{idx}(a)] = D_{\text{cur}}[\text{idx}(b)]\} \quad (1)$$

consists of a vertex-disjoint union of directed cycles that cover all vertices in S . This means that S is a set of patients that can all get their preferred doctor if we allow for simultaneous exchanges following cycles.

3.1.3. Optimization criterion

Next we can start defining variants of feasible solutions where we want to maximize some *score*. Examples of such a score could be to find a feasible solution with many patients or something based on priority. We might want a solution that exchanges as many patients as possible, exchanges patients with the highest priority or some mix of these.

First we define our general ordering

Definition 3.1.3.1 (Optimization criterion): An ordering $\succ$ over feasible solutions. A solution is optimal if it is maximal under $\succ$.

Build on this to create our three main variants of scoring functions.

3.1.3.1. Lexicographic maximization by priority

Definition 3.1.3.1.1 (Characteristic vector): Let patients be ordered by priority: $p_1 \geq p_2 \geq \dots \geq p_n$ where p_1 has the highest priority. Given a feasible solution $S \subseteq P$, the **characteristic vector** is:

$$\chi(S) = (b_1, b_2, \dots, b_n) \in \{0, 1\}^n, \quad b_i = \begin{cases} 1 & \text{if } p_i \in S \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

Definition 3.1.3.1.2 (Lexicographic ordering by priority):

$$S \succ_{\text{lex}} S' \quad \text{iff} \quad \chi(S) \text{ is lexicographically greater than } \chi(S') \quad (3)$$

That is, at the first index i where S and S' differ, $\chi(S)_i = 1$ and $\chi(S')_i = 0$.

Intuitively, $\chi(S)$ is a binary number whose most significant non-zero bit corresponds to the highest-priority patient that gets helped. The optimal solution is then the one that maximises this binary number: always satisfying the highest-priority patient it can, then subject to that the next highest, and so on.

Example (Ordering two solutions lexicographically by priority):

We use the example graph in Figure 4 as the graph to create solutions from.

Given solutions

1. $S = \{P_4, P_5\}$ represents the subgraph G_S containing the cycle $P_4 \rightarrow P_5 \rightarrow P_4$
2. $S' = \{P_1, P_2, P_3\}$ represents the subgraph $G_{S'}$ containing the cycle $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_1$

For simplicity's sake we say that $R(P_i) = i$, so P_5 has the highest priority.

Ordering all five patients by priority gives $p^{(1)} = P_5, p^{(2)} = P_4, \dots, p^{(5)} = P_1$. Then:

$$\chi(S) = (1, 1, 0, 0, 0) \quad \chi(S') = (0, 0, 1, 1, 1) \quad (4)$$

At the first differing position ($i = 1$), $\chi(S)_1 = 1 > 0 = \chi(S')_1$, so $S \succ_{\text{lex}} S'$.

So for our first variant we want to find the maximal solution under the $\succ_{\text{lex}}$ ordering. This means always prioritizing patients with highest priority.

3.1.3.2. Maximizing cardinality

Another natural variant is to find a solution with the largest cardinality, e.g. a solution that contains the most patients.

Definition 3.1.3.2.1 (Ordering by cardinality):

$$S \succ_{\text{size}} S' \quad \text{iff} \quad |S| > |S'| \quad (5)$$

This optimal solution will then be one that exchanges the most patients.

Example (Ordering two solutions by cardinality):

Using the same solutions as in the previous example:

1. $S = \{P_4, P_5\}$ represents the subgraph G_S containing the cycle $P_4 \rightarrow P_5 \rightarrow P_4$
2. $S' = \{P_1, P_2, P_3\}$ represents the subgraph $G_{\{S'\}}$ containing the cycle $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_1$

Under $\succ_{\text{size}} : |S| = 2$ and $|S'| = 3$, so $S' \succ_{\text{size}} S$. However the solution is $\{P_1, P_2, P_3, P_4, P_5\}$, the union of both cycles.

Notice that the same two sets are reversibly ordered under the two criteria: $S \succ_{\text{lex}} S'$ (because S contains the highest-priority patient P_5) but $S' \succ_{\text{size}} S$ (because S' has more patients).




Figure 4: An example graph G .

3.1.3.3. Maximizing utility

Finally we formulate an ordering that is a combination of $\succ_{\text{lex}}$ and $\succ_{\text{size}}$. We define the total utility to be the sum of the priorities of the patients in S .

We recall that $R : P \rightarrow \mathbb{N}$ is the function mapping a patient to a priority.

Definition 3.1.3.3.1 (Total utility function):

$$U(S) = \sum_{p \in S} R(p) \tag{6}$$

Definition 3.1.3.3.2 (Ordering by total utility):

$$S \succ_{\text{util}} S' \text{ iff } U(S) > U(S') \quad (7)$$

This means that even though solution S' contains some high priority patients, if S contains enough lower priority patients then S can still be a higher ranked solution under $\succ_{\text{util}}$.

Example (Ordering two solutions by total utility):

We use the following solutions from Figure 5:

1. $S = \{1, 3, 4\}$ representing the cycle $P_1 \rightarrow P_3 \rightarrow P_4 \rightarrow P_1$
2. $S' = \{5, 2\}$ representing the cycle $P_5 \rightarrow P_2 \rightarrow P_5$

This gives the following total utilities:

$$U(S) = 1 + 3 + 4 = 8$$

$$U(S') = 2 + 5 = 7$$

It follows that $S \succ_{\text{util}} S'$. Here we see that the solution with more patients with lower maximum priority has a higher score than another with greater maximum priority.




Figure 5: An example graph G.

3.1.4. Priority function

The priority function has quite a large effect on the switches we end up taking. As the priority of a patient decides if that patient will be in the solution or not for most of the algorithms, except for the exact algorithm maximizing cardinality $\succ_{\text{size}}$. This is why its important to discuss what effects the function can have on solutions when we are maximizing the ordering $\succ_{\text{util}}$.

If we make the priority function exponential such that $R(a) = 2^a$ then using the ordering $\succ_{\text{util}}$ becomes equal to $\succ_{\text{lex}}$. While this ordering is in terms maybe the most fair it might not always be the best for the collective good, since a high priority patient might block a lot of lower

priority patients from switching. This way we might want a priority function like $R(a) = \text{Days patient } a \text{ has been waiting for a switch}$. This way if we have the choice between patient P_i who has waited for 30 days $R(P_i) = 30$ and 4 patients who have waited 40 days in total an algorithm maximizing $\succ_{\text{util}}$ will choose the 4 patients.

We can adjust the priority function to prioritize higher priority by making it exponential but still under the 2^a . When we use days waited we can make $R(a) = 1.1^{\text{Days patient } a \text{ has been waiting}}$. Adjusting this closer to 2 makes higher priority patients more prioritized and vice versa for lowering it.

4. Implementation

In this chapter we look at each algorithm we have implemented, how we implemented them, why and runtime analysis. We have in total implemented 5 algorithms, TTC, Greedy DFS, Exact Cardinality, Exact Priority and a specialized case of Exact Priority. For the exact algorithms we also give, or refer to existing, proofs for why they are exact.

4.1. Different graph representations used

In the algorithms we use different graph representations, so we start by defining the different representations of *The GP allocation problem*.

4.1.1. Patient and Doctor graph

First we have the easiest graph containing both patients and doctors as vertices. The edges are directed from a patient to a doctor and from a doctor to a patient. Edges never go patient $\rightarrow$ patient or doctor $\rightarrow$ doctor. An edge patient $\rightarrow$ doctor symbolizes that the patient has that doctor as its preferred doctor, and doctor $\rightarrow$ patient symbolises that the doctor currently has that patient as one of his or hers current patients. Observe that this graph is bipartite as we have patients on one side and doctors on the other. Now the formal definition of our graph. Let $I = \{0, \dots, |P| - 1\}$. Then:

$$G = (V, E), \quad V = P \cup D, \quad E = \{(p_i, D_{\text{pref}[i]}) \mid i \in I\} \cup \{(D_{\text{cur}[i]}, p_i) \mid i \in I\} \quad (8)$$

4.1.2. Doctor graph collapsed edges

In this weighted graph we condense the problem to only have doctors as nodes and edges between doctors. An edge doctor $a \rightarrow$ doctor b symbolises that there exists a patient that wants to switch from doctor a to doctor b , or that the patient currently has doctor a and has doctor b as preferred. We define our graph as:

$$\begin{aligned} G &= (V, E), \quad V = D \\ E &= \{(D_{\text{cur}[i]}, D_{\text{pref}[i]}) \mid i \in I\} \\ w(a, b) &= |\{i \in I \mid D_{\text{cur}[i]} = a \wedge D_{\text{pref}[i]} = b\}| \end{aligned} \quad (9)$$

4.2. Greedy DFS

We first consider a greedy approach inspired by the Top Trading Cycles implementation of Huitfeldt et al. Their algorithm preserves TTC properties at each round by restricting each doctor node to a single outgoing edge. We relaxed this constraint, allowing each doctor to maintain outgoing edges to all of its current patients simultaneously, and resolved ties by patient priority.

We are given the lists of patients and doctors P, D , the assignment arrays $D_{\text{cur}}, D_{\text{pref}}$, and a priority function $R : P \rightarrow \mathbb{N}$, where a higher number means higher priority. We model the problem as a bipartite directed graph over patients and doctors. Let $I = \{0, \dots, |P| - 1\}$. Then:

$$G = (V, E), \quad V = P \cup D, \quad E = \{(p_i, D_{\text{pref}[i]}) \mid i \in I\} \cup \{(D_{\text{cur}[i]}, p_i) \mid i \in I\}$$

Each patient p_i has an outgoing edge to their preferred doctor, and each doctor has outgoing edges to all of its currently registered patients. A directed cycle in G necessarily alternates between patient and doctor nodes and corresponds to a valid exchange: each patient in the cycle moves to their preferred doctor, and each doctor loses exactly one patient while gaining one.

The algorithm processes patients in decreasing priority order. For each patient, a DFS attempts to find a cycle through that patient; when the DFS reaches a doctor node with multiple candidate patients, it always explores the highest-priority one first.

```

1 let resolved_patients = []
2 let p_prio = Sorted list of patients by priority
3 let wants_to_switch = [true] * |P|
4 for each  $p \in p\_prio$  do
5   if let cycle = dfs_find_cycle( $G, R, p$ ) then
6     for each  $p \in cycle$  do
7       let  $i = \text{idx}(p)$ 
8       wants_to_switch[ $i$ ] = false
9       resolved_patients.push( $p$ )
10    end
11  end
12 end
13 return resolved_patients

```

The greedy rule ensures that high-priority patients are preferentially included in cycles, but it does not guarantee a globally optimal solution. The following example illustrates how the greedy choice can be locally motivated yet globally suboptimal.




Figure 6: Example graph G where algorithm would choose suboptimal solution. p_x means patient x and has priority x .

In Figure 6 each patient p_x has priority x . The DFS begins at p_4 (highest priority), follows the edge to d_2 , and there chooses p_2 over p_1 since $R(p_2) > R(p_1)$. The resulting cycle $p_4 \rightarrow d_2 \rightarrow p_2 \rightarrow d_1 \rightarrow p_4$ is committed, leaving p_1 and p_3 unsatisfied with no further cycles remaining.

Had the DFS chosen $p1$ at $d2$ instead, it would have found the longer cycle $p4 \rightarrow d2 \rightarrow p1 \rightarrow d3 \rightarrow p3 \rightarrow d1 \rightarrow p4$, satisfying three patients. The greedy choice at $d2$ was locally motivated by priority but globally suboptimal. This motivates the exact algorithms in the following sections.

4.3. Exact algorithm for maximizing total switches

NB: HER OG NESTE ER DET MYE AI, NOE JEG VIL HØRE DERES MENING OM. NOE AV DETTE SYNES JEG ER SKREVET BRA, ANDRE LITT USIKKER. INKLUDERT KILDENE ER IKKE RIKTIG

This algorithm applies the classical cycle canceling technique from network flow theory to the GP-switching problem, finding the maximum-cardinality feasible solution in polynomial time [10, §9.6]. The key insight is that the problem reduces to a maximum integer circulation, for which efficient algorithms are known.

4.3.1. Graph structure

Rather than working with the bipartite patient-doctor graph, we compress to a weighted directed graph over doctors only. Patients who want the same switch are interchangeable: all that matters is how many can be routed. We therefore aggregate them into edge weights:

$$\begin{aligned} G &= (V, E), \quad V = D \\ E &= \left\{ (D_{\text{cur}[i]}, D_{\text{pref}[i]}) \mid i \in \llbracket P \rrbracket \right\} \\ w(a, b) &= |\{i \in \llbracket P \rrbracket \mid D_{\text{cur}[i]} = a \wedge D_{\text{pref}[i]} = b\}| \end{aligned} \tag{10}$$

A cycle $d_1 \rightarrow d_2 \rightarrow \dots \rightarrow d_k \rightarrow d_1$ carrying f units of flow corresponds to f patients rotating around the cycle: each moves from their current doctor to the next in the cycle. This compression reduces the graph from $O(|P| + |D|)$ nodes to $O(|D|)$ nodes, which is significant when many patients share the same switch request.

Applying this transformation to Figure 6 gives the doctor graph in Figure 7, where we can still identify the short cycle $d1 \rightarrow d2$ and the longer cycle $d1 \rightarrow d2 \rightarrow d3$.




Figure 7: Graph from Figure 6 but in doctor graph structure.

4.3.2. Algorithm

We want to find, for each edge $e \in E$, a non-negative integer $f(e)$ representing how many of the $w(e)$ patients on that edge actually switch. Two conditions make such an assignment a valid set of simultaneous exchanges:

- **Capacity:** $0 \leq f(e) \leq w(e)$ for all $e \in E$

- **Flow conservation:** for every doctor $d \in D$: $\sum_{(v,d) \in E} f(v,d) = \sum_{(d,v) \in E} f(d,v)$

Flow conservation is exactly what forces the assignment to decompose into cycles: every doctor who loses a patient must gain one. The problem is therefore:

Problem: Maximum Switch Circulation

Input: A directed graph $G = (D, E, w)$ where $w : E \rightarrow \mathbb{N}$

Output: An integer function $f : E \rightarrow \mathbb{N}_0$ maximising $\sum_{e \in E} f(e)$ subject to:

- $0 \leq f(e) \leq w(e)$ for all $e \in E$
- $\sum_{(v,u) \in E} f(v,u) = \sum_{(u,v) \in E} f(u,v)$ for all $u \in D$

This is a **maximum integer circulation** problem, solvable in polynomial time. We solve it using the **successive positive cycle** method.

Given a current flow f , the **residual graph** G_f is built by replacing each original edge (u, v) with:

- a **forward residual** arc (u, v) with capacity $w(u, v) - f(u, v)$ and cost $+1$
- a **backward residual** arc (v, u) with capacity $f(u, v)$ and cost -1

A directed cycle in G_f is **positive** if the sum of its arc costs is positive, meaning it contains more forward than backward arcs, so augmenting along it strictly increases total flow. The optimality condition for maximum circulations states that f is optimal if and only if G_f contains no positive cycle.

The outer loop finds and augments positive cycles until none remain:

```

1 for each  $e \in E$  do  $f(e) \leftarrow 0$  end
2 while FindPositiveCycle( $G_f$ ) returns a cycle  $C$  do
3    $b \leftarrow \min_{e \in C} r_f(e)$   $\triangleright$  bottleneck residual capacity
4   for each  $(u, v) \in C$  do
5      $f(u, v) \leftarrow f(u, v) + b$ ; update residual capacities in  $G_f$ 
6   end
7 end
8 return  $f$ 

```

Positive cycles are detected using a variant of Bellman–Ford (SPFA). All nodes are seeded with distance 0, simulating a virtual super-source at zero cost so every cycle is reachable. Edges are relaxed greedily in the maximising direction; a node relaxed $|D|$ or more times must lie on a positive cycle.

```

1 for each  $v \in D$  do  $\text{dist}[v] \leftarrow 0$ ;  $\text{pred}[v] \leftarrow \perp$ ; enqueue  $v$  end
2 while queue not empty do
3    $u \leftarrow \text{dequeue}$ 

```

```

4  for each  $(u, v)$  in  $G_f$  with  $r_f(u, v) > 0$  do
5      if  $\text{dist}[u] + 1 > \text{dist}[v]$  then
6           $\text{dist}[v] \leftarrow \text{dist}[u] + 1$ ;  $\text{pred}[v] \leftarrow u$ 
7          if  $v$  has been enqueued  $\geq |D|$  times then
8              Walk back  $|D|$  steps from  $v$  along  $\text{pred}$  to find node  $c$  on the cycle
9              Collect and return the cycle starting and ending at  $c$ 
10         end
11         Enqueue  $v$  if not already in queue
12     end
13 end
14 end
15 return  $\emptyset$        $\triangleright$  no positive cycle;  $f$  is optimal

```

Theorem 4.3.2.1 (Optimality of MaxSwitchCirculation): The algorithm returns a flow f achieving the maximum possible value $\sum_{e \in E} f(e)$, equal to the maximum number of patients that can simultaneously switch to their preferred doctor.

Theorem 7: Optimality of MaxSwitchCirculation

Proof: The algorithm terminates when FindPositiveCycle returns $\emptyset$, meaning G_f contains no positive-cost directed cycle. By the cycle optimality criterion for maximum circulations: a feasible integer circulation f is optimal if and only if its residual graph G_f contains no positive-cost cycle. Since each forward arc carries cost $+1$ and each backward arc cost -1 , any positive-cost cycle in G_f would strictly increase $\sum f(e)$; the absence of such cycles certifies that f is maximum. $\square$

Termination and complexity. Each call to FindPositiveCycle runs SPFA over $|D|$ nodes and at most $|D|^2$ edges in $O(|D|^3)$ time. Each augmentation round increases total flow by at least 1, so the number of rounds is at most $W = \sum_{e \in E} w(e)$, the total number of patients wanting to switch, giving a worst-case bound of $O(W \cdot |D|^3)$. In practice each round pushes a bottleneck of $b \geq 1$ units, so the number of distinct rounds is bounded by the number of edges that become saturated, at most $|E| \leq |D|^2$. This gives the graph-size bound $O(|D|^5)$, independent of the patient count.

4.4. Exact algorithm for maximizing priority

This algorithm finds the lexicographically maximal feasible solution under $\succ_{\text{lex}}$, as defined in Section 3. It builds on the same residual-graph framework but processes patients in priority order, greedily committing each one as soon as a cycle is found. The correctness of this greedy strategy follows from the theorem below.

Theorem 4.4.1 (Greedy choice property): Let p^* be the highest-priority patient that belongs to some feasible solution. Then every lexicographically maximal solution contains p^* .

Theorem 8: Greedy choice property

Proof: Suppose S^* is a lexicographically maximal feasible solution with $p^* \notin S^*$. By hypothesis there exists a feasible solution S' with $p^* \in S'$. Let k be the index of p^* in the priority ordering $p^{(1)} \succ p^{(2)} \succ \dots \succ p^{(n)}$. Every patient with index $i < k$ belongs to no feasible solution by the choice of p^* , so $\chi(S^*)_i = \chi(S')_i = 0$ for all $i < k$. At index k : $\chi(S')_k = 1 > 0 = \chi(S^*)_k$. Therefore $S' \succ_{\text{lex}} S^*$, contradicting the maximality of S^* . $\square$

The algorithm applies this observation iteratively: process patients from highest to lowest priority and commit each patient to a cycle as soon as one is found.

We use the same doctor graph as before: doctors are nodes, and each patient p with current doctor u and preferred doctor v is a directed arc $a_p = u \rightarrow v$ with capacity 1. Each arc also has a corresponding **backward arc** $\overline{a_p} = v \rightarrow u$ with initial capacity 0.

Pruning. Before processing, we compute the strongly connected components (SCCs) of the graph. Any patient whose current and preferred doctor lie in different SCCs, or in a trivial SCC of size 1, can never be part of any directed cycle and is immediately discarded. This avoids unnecessary BFS calls and keeps the residual graph clean throughout execution.

The algorithm processes the remaining patients from highest to lowest priority. For each patient p (arc $u \rightarrow v$), exactly one of two cases applies:

- **Case 1 (primary arc available, $\text{cap}(a_p) > 0$):** p has not yet been committed. We run BFS from v to u in the current residual graph. If a path Π exists, then Π together with a_p forms a cycle, and we commit p :
 - The primary arc a_p is consumed **permanently**: its backward arc capacity stays 0, so no future patient can traverse it in reverse and undo p 's switch.
 - Each routing arc a in Π is consumed normally: $\text{cap}(a)$ decreases by 1 and $\text{cap}(\overline{a})$ increases by 1, leaving a residual that lower-priority patients may use to reroute.
 - If no path exists, p is skipped.
- **Case 2 (primary arc consumed, residual active, $\text{cap}(a_p) = 0$ and $\text{cap}(\overline{a_p}) > 0$):** p 's arc was already consumed as a **routing arc** in an earlier (higher-priority) patient's cycle, which created the residual $\overline{a_p}$. We now **solidify** by setting $\text{cap}(\overline{a_p}) \leftarrow 0$, preventing any lower-priority patient from traversing this residual to undo p 's assignment.

- 1 Compute SCCs of G ; discard patients whose arc spans different or trivial SCCs
- 2 $E_{\text{sort}} \leftarrow$ remaining patients sorted by priority descending
- 3 **for each** patient p with arc $a_p = (u \rightarrow v)$ in E_{sort} **do**
- 4 **if** $\text{cap}(a_p) > 0$ **then**
- 5 **if** $\text{BfsPath}(v, u)$ in residual graph returns path Π **then**

```

6   |   |   cap( $a_p$ )  $\leftarrow$  0                                ▷ consume primary permanently; no residual
7   |   |   for each arc  $a \in \Pi$  do
8   |   |   |   cap( $a$ )  $\leftarrow$  cap( $a$ ) - 1; cap( $\bar{a}$ )  $\leftarrow$  cap( $\bar{a}$ ) + 1
9   |   |   end
10  |   |   Add  $p$  to  $S$ 
11  |   |   end
12  |   |   else if cap( $\bar{a}_p$ ) > 0 then
13  |   |   |   cap( $\bar{a}_p$ )  $\leftarrow$  0                                ▷ solidify; lock in routing commitment
14  |   |   |   Add  $p$  to  $S$ 
15  |   |   end
16  |   end
17  return  $S$ 

```

Why the primary arc has no residual: This is the core invariant. Once patient p is committed as the beneficiary of a cycle, no future patient should be able to undo it. By not creating a residual for a_p , no BFS can ever traverse backwards through p 's switch.

Why routing arcs do have residuals: A lower-priority patient may be able to “take over” a routing role. For example, if p_1 (high priority) uses p_3 's arc as routing, a later patient p_2 may form a cycle that routes through p_3 's arc differently, replacing it. This is valid: only the final destination of p_1 (whose primary arc is irrevocable) is fixed.

Theorem 4.4.2 (Correctness): The algorithm returns the lexicographically maximal feasible solution.

Theorem 9: Correctness

Proof: By induction on the priority ordering. The Greedy Choice Theorem establishes the base case: the highest-priority feasible patient p^* must appear in every lex-max solution, and BFS correctly identifies whether p^* is feasible, since a path from v to u exists in G if and only if p^* lies on some directed cycle.

For the inductive step, assume all higher-priority patients have been correctly committed. The residual graph reflects these commitments: primary arcs are permanently consumed and cannot be undone, while routing residuals remain available. We claim a path from v to u exists in the current residual graph if and only if p can be added to a feasible solution extending all committed patients. This follows from the augmenting-path argument: any feasible cycle through p either avoids all committed arcs and is directly reachable in G_f , or shares some routing arcs whose residuals permit an equivalent rerouting. The solidify step in Case 2 maintains the invariant: once a patient's arc is confirmed as routing for a higher-priority cycle, its residual is removed so no lower-priority patient can undo the commitment. $\square$

Complexity. Sorting patients requires $O(|P| \log |P|)$. The initial SCC computation runs in $O(|D| + |E|)$. Each patient requires at most one BFS over the residual graph in $O(|D| + |E|)$ time. Processing all $|P|$ patients therefore takes $O(|P| \cdot (|D| + |E|))$. Since $|E| \leq |D|^2$, the overall complexity is $O(|P| \cdot |D|^2)$.

4.5. Metaheuristics

5. Experimental Results

5.1. Experimental Setup

5.2. Performance Comparison

5.3. Impact of Priority Strategies

6. Conclusion

6.1. Summary

6.2. Future Work

Bibliography

- [1] Legelisten, “Med rett til å bytte – Ventetider for å bytte fastlege i Norge.” Accessed: Apr. 08, 2026. [Online]. Available: <https://legelisten.s3.eu-west-1.amazonaws.com/blog-files/Med+rett+til+%C3%A5+bytte++Ventetider+for+%C3%A5+bytte+fastlege+i+Norge.pdf>[◦]
- [2] Helfo, “Fastlegestatistikk.” Accessed: Apr. 08, 2026. [Online]. Available: <https://www.helfo.no/Fastlegeordninga/fastlegestatistikk>[◦]
- [3] Statistisk sentralbyrå, “Fastlegelister og fastlegekonsultasjoner, etter år og region (12005).” Accessed: Apr. 08, 2026. [Online]. Available: <https://www.ssb.no/statbank/table/12005>[◦]
- [4] L. Shapley and H. Scarf, “On cores and indivisibility,” *Journal of Mathematical Economics*, vol. 1, no. 1, pp. 23–37, 1974, doi: [https://doi.org/10.1016/0304-4068\(74\)90033-0](https://doi.org/10.1016/0304-4068(74)90033-0)[◦].
- [5] D. J. Abraham, A. Blum, and T. Sandholm, “Clearing algorithms for barter exchange markets: enabling nationwide kidney exchanges,” pp. 295–304, 2007, doi: [10.1145/1250910.1250954](https://doi.org/10.1145/1250910.1250954)[◦].
- [6] Wikipedia contributors, “Optimal kidney exchange — Wikipedia, The Free Encyclopedia.” Accessed: May 05, 2026. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Optimal_kidney_exchange&oldid=1351916222[◦]
- [7] D. Gale and L. S. Shapley, “College Admissions and the Stability of Marriage,” *The American Mathematical Monthly*, vol. 69, no. 1, pp. 9–15, 1962, Accessed: May 06, 2026. [Online]. Available: <http://www.jstor.org/stable/2312726>[◦]
- [8] Wikipedia contributors, “Top trading cycle — Wikipedia, The Free Encyclopedia.” Accessed: May 08, 2026. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Top_trading_cycle&oldid=1344910705[◦]
- [9] A. E. Roth, “Incentive compatibility in a market with indivisible goods,” *Economics Letters*, vol. 9, no. 2, pp. 127–132, 1982, doi: [https://doi.org/10.1016/0165-1765\(82\)90003-9](https://doi.org/10.1016/0165-1765(82)90003-9)[◦].
- [10] R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, *Network Flows: Theory, Algorithms, and Applications*. Prentice Hall, 1993.

Index of Figures

Figure 1	An example graph for the kidney exchange problem.	8
Figure 2	An example Housing Market for TTC.	9
Figure 3	An example Housing Market for TTC, A's $\text{Top}(i)$ is itself.	9
Figure 4	An example graph G	12
Figure 5	An example graph G	13
Figure 6	Example graph G where algorithm would choose suboptimal solution. px means patient x and has priority x	16
Figure 7	Graph from Figure 6 but in doctor graph structure.	17