

Top Trading Cycles in a practical setting

Author: Lars Møen Haukland

Supervisor: Fredrik Manne



UNIVERSITETET I BERGEN
Fakultet for naturvitenskap og teknologi

June, 2026

Abstract

This thesis explores the Top Trading Cycles algorithm in the practical setting of the allocations of GPs, General Practitioner, in Norway for people wanting to switch or get a GP. It explores how different focus of *good* results can change both the runtime and results, and how to make the algorithm run as efficiently as possible. This thesis proposes an implementation of the Top Trading Cycles that can reduce the waitlists for GPs in Norway by over 80%.

Acknowledgements

I would like to thank my supervisor Fredrik Manne for his guidance and support throughout this project. I would also like to thank Juni Weisteen Bjerde for input and help during the project. In addition Juni and Mathias Bertntert together theorised the exact algorithms so that I could implement them.

Contents

1. Introduction	5
1.1. Thesis Outline	5
2. Background	7
2.1. Related problems	7
2.1.1. One-to-One assignment problems	7
2.1.1.1. Housing market	7
2.1.1.2. Kidney exchange	7
2.1.2. Many-to-One assignment problems	8
2.1.2.1. The College Admissions problem	8
2.2. Top Trading Cycles	8
2.2.1. Classical TTC	8
2.2.2. GP assignment problem TTC	10
2.3. Cycle cancelling	10
2.3.1. Minimum cost maximum circulation problem	10
2.3.2.	11
2.3.3. Algorithm	11
2.3.4. Runtime	11
3. Problem Formalization	13
3.1. The GP allocation problem	13
3.1.1. Input	13
3.1.2. Feasible solution	13
3.1.3. Optimization criterion	13
3.1.3.1. Lexicographic maximization by priority	14
3.1.3.2. Maximizing cardinality	14
3.1.3.3. Maximizing utility	15
3.1.4. Priority function	16
4. Implementation	18
4.1. Different graph representations used	18
4.1.1. Patient and Doctor graph	18
4.1.2. Doctor graph collapsed edges	18
4.1.3. Doctor graph priority weighted	18
4.2. Greedy DFS	19
5. Experimental Results	21
5.1. Experimental Setup	21
5.2. Performance Comparison	21
5.3. Impact of Priority Strategies	21
6. Conclusion	22
6.1. Summary	22
6.2. Future Work	22
Bibliography	23
Index of Figures	24

1. Introduction

The Norwegian general practitioner (GP) system, started in 2001, has aim to serve the Norwegian population with a stable GP. A GP can follow his or hers patients over time, creating a trustful and safe relationship [1]. Today this system works as intended, and GPs handle most of the populations needs. As of December 2025 the number of GPs in Norway is 5720 and the total number of citizens in the Norwegian GP system is 5.6 million. The amount of patients each GP has varies but on average each GP has about a thousand patients [2].

One inefficiency in the GP system is when a person wants to switch from one GP to another. If the GP they want to switch to does not have any more capacity for new patients, the person has to sign up in a queue and wait for a free slot to open up with that GP. In December the number of open GP lists was 1340, indicating that 77% of all GPs had no extra capacity [3]. In a more densely populated city such as Bergen, the numbers are even worse. In December 2025 Bergen had 268 GPs and only 7 GPs that had available capacity for new patients. That's 97% of the GPs in Bergen had no extra capacity [3].

When a person wants to switch to a GP that has no excess capacity they must sign up and then wait in line. In December 2025 there were about 300,000 people waiting for GPs in Norway [2]. As there are only 5720 doctors each doctor has on average more than 50 people on his or hers waiting list. Since many are switching from a GP it follows that it is likely that there are cycles, in that a person P_A wants doctor D_A which currently has P_B as a patient, but P_B wants to switch to doctor D_B who currently has P_A as a patient. In this case patient P_A ends up waiting for an open slot with P_B 's doctor and vice versa, but logically instead of waiting P_A and P_B could just swap doctors. This can be generalized to longer "waiting cycles" that could all be serviced by allowing for simultaneous switching.

In this thesis, we study this problem and come up with solutions for how one can help people who are waiting in line so that they can get their preferred choice. In particular we propose that we can use cycle detecting algorithms, to find optimal switches between patients and massively reduce the number of people on waiting lists by using these algorithms periodically. Our algorithms are based on existing cycle cancelling and TTC algorithms, and focus on two measures of fairness:

- Switching as many patients as possible
- Switching as high priority patients as possible

We have developed and tested algorithms for these cases.

In addition we compare these optimal algorithms with the existing TTC algorithm and a greedy algorithm, comparing gain of good vs runtime.

1.1. Thesis Outline

The remainder of this thesis is organized as follows:

- **Section 2** provides background on the Top Trading Cycles mechanism, cycle cancelling algorithms and how these methods have been used in similar problems.
- **Section 3** formally defines the computational problems arising from TTC as variants of cycle cover problems on directed graphs.

- **Section 4** describes our Rust implementation and the algorithmic strategies employed.
- **Section 5** presents experimental results comparing different priority strategies.
- **Section 6** concludes and discusses future work.

2. Background

Problems like the GP allocation problem have been studied for quite some time with a lot of variations. In addition there are different perspectives on how to solve these problems. While in economics they focus on the fairness of an assignment, trying to make a stable and strategy proof assignment. In this chapter we describe similar problems for the GP allocation problem, describe the Top Trading Cycles algorithm (TTC) and Huitfeldt et.al implementation of TTC for the GP allocation problem. Finally we describe cycle cancelling, an optimization technique that our exact algorithms are based on.

2.1. Related problems

There are variations of assignment problems. Usually in assignment problems we mention agents and objects and agents want objects, in the GP allocation problem the patients would be agents and doctors would be objects. The typical assignment problems are One-to-One assignment problems, meaning one agent gets one object and an object can only be “owned” by one agent. In addition we also have Many-to-One assignment problems, like the GP allocation problem, one agent has one object but one object can be “owned” by more than one agent. It might be misleading to say that our patients “own” doctors but this makes more sense in other problems that we will describe, owning a doctor in our case means having that doctor as a current doctor.

2.1.1. One-to-One assignment problems

2.1.1.1. Housing market

The Housing Market model introduced in Shapley and Scarf [4] describes a model with traders (agents) and indivisible goods (objects). These goods can be houses, making it a housing market. Each agent already has one house but may want to switch. Every agent has a preference list ordering every house in order of preference, Shapley and Scarf combine all these preference lists to make a preference matrix [4].

It is in this article from Shapley and Scarf [4] that they also introduce the Top Trading Cycles algorithm and describe how it is fitting for models like the Housing Market. We also see similarities to the GP allocation problem in that we have agents already owning an object but wanting to switch. Two differences is the One-to-One since a doctor can have multiple patients compared to that one house can only be owned by one agent, and that in the Housing market each agent has a complete preference list while in the GP assignment problem each agent has only one preferred doctor.

2.1.1.2. Kidney exchange

The Kidney Exchange problem describes an issue when trying to find compatible kidney donors. It describes the case where we have pairs of patients and donors, but the patient is incompatible with the donors kidney. We then need to find some way to swap around the donors to compatible patients. Either with pairwise swapping or with larger cycles. If we visualize the problem as a directed graph, each vertex is a patient and donor pair. An edge from Pair A to Pair B means that Donor A is compatible with Patient B. Now cycles can form as in Figure 1 we have a cycle of length 3 $A \rightarrow B \rightarrow C$.




Figure 1: An example graph for the kidney exchange problem.

Abraham, Blum and Sandholm [5] show that, with unbounded cycle length, a maximum-cardinality and maximum-weight exchange can be found in polynomial time [6]. An issue often arising with unrestricted cycle length is if a donor suddenly refuses to donate, if this donor refuses after their patient has received a kidney then a patient is left without a donor and cannot exchange later. Because of this the problem is often considered with a restricted cycle length to allow all operations to be executed at the same time. This way no donor can refuse in the middle. For each pair in a cycle you need 2 operations, for a cycle of length k you need $2k$ operations at the same time. Finding a maximum-cardinality exchange with cycles of length at most k , for any fixed $k \geq 3$, is an NP-hard computational problem [5], [6].

The Kidney Exchange problem and the GP allocation problem have much in common. In both problems agents already have an object, but want to switch for another. We need to find cycles to exchange the objects in such a manner that most amount of people are happy. The key difference is that only one donor can be “owned” by one patient while in the GP allocation problem a doctor can be “owned” by multiple patients.

2.1.2. Many-to-One assignment problems

2.1.2.1. The College Admissions problem

In the College Admissions problem, introduced by Gale and Shapley, describes the problem of assigning applicants among colleges. Each college has a certain quote, a maximum number of applicants to admit and each applicant ranks the colleges in order of preference. An applicant can skip colleges he or she would never accept under any circumstances [7].

The problem is then to find a stable assignment meaning that there does not exist a case where there are two applicants α and β who are assigned to colleges A and B , respectively, although β prefers A to B and A prefers β to α . Gale and Shapley introduce an algorithm called deferred-acceptance (DA) to find the optimal stable assignment and prove that this runs in polynomial time and finds the optimal stable assignment [7].

There are close similarities between the College Admissions problem and the GP assignment problem. One difference is that patients, equivalent to the applicants, only prefer one doctor. Another difference is that patients already have existing doctors, making the problem different since applicants do not already have a college they want to switch from.

2.2. Top Trading Cycles

2.2.1. Classical TTC

Top Trading Cycles (TTC), developed by Gale and published by Shapley and Scarf [4], is an algorithm to find a re-allocation of goods without using money. As mentioned it was

introduced in an article together with the Housing Market problem. The algorithm finds a re-allocation of houses to traders, such that all mutually-beneficial exchanges have been realized [8]. For the Housing Market problem the TTC algorithm does as follows [8]:

1. Ask each agent to indicate his “top” (most preferred) house.
2. Draw an arrow from each agent i to the agent, denoted $\text{Top}(i)$, who holds the top house of i .
3. Find a cycle (guaranteed to exist) and execute the trades in that cycle. Remove all involved agents from the graph.
4. If there are remaining agents, go back to step 1.

Note that a cycle is guaranteed to exist if there are agents still left. This is because of the pidgeonhole principle, since each agent has one outgoing edge we have to have a cycle. The cycle can also be of length 1 in which case an agent’s top house is its own, then we treat it as a normal cycle and remove the agent from the graph. For this classical TTC implementation to work we need to remember that agents have a strict preferences. Each agent has a complete list of all houses ranked in order of preference. Because of this as houses get re-allocated agents are bound to either find a trade cycle with others or end up having their $\text{Top}(i)$ house as their own.

Consider the example of 3 agents, A , B and C , each has a house and have a strict preference list. Lets go through how TTC would handle this:

The preference lists are

- $A = [B, C, A]$
- $B = [C, A, B]$
- $C = [B, C, A]$

If we make the graph where each agent point to their top we get Figure 2.




Figure 2: An example Housing Market for TTC.

We have the cycle $B \rightarrow C$ so we execute the trades, remove agents from the graph and create the edges again. The last remaining node is A and as B and C are not in the graph, $\text{Top}(i)$ of A is now A .




Figure 3: An example Housing Market for TTC, A 's $\text{Top}(i)$ is itself.

We execute this trade and are left with no more agents, making TTC terminate.

Roth proved that TTC is strategy proof, meaning all agents are motivated to prefer their true preferences [9]. TTC has also satisfies properties of Individual rationality, that an agent is at least as well off by participating, and Pareto efficiency which is that the objects are allocated such that no agent can improve without worsening another agent.

Now using TTC for the GP assignment problem is challenging as patients do not have complete strict preference lists, they only have one preferred doctor, and a doctor can be “owned” by multiple patients. This makes the $\text{Top}(i)$ give out multiple patients, as the doctor that patient i prefers is “owned” by multiple patients. How should we choose between these? This is what Huitfeldt et. al has written about and how their implementation satisfies the same properties of the classical TTC. Later we run the TTC made for the GP assignment problem and compare with our algorithms. First we explain it and how it is different from normal TTC.

2.2.2. GP assignment problem TTC

In the paper from Huitfeldt et al. they tackle a more complex dataset as they use real patient and doctor data from Norway. In their model they have doctors with available capacity and patients switching to these GPs can be allocated that GP without an exchange. This is why in each of the two TTC implementations they first run a waitlist algorithm that goes through all GP with available capacity and assigns the highest priority to that GPs panels until there are no more patients waiting on GPs with available capacity [10]. Following we explain Huitfeldt et al. TTC implementation without regard to capacities to doctors as that is how we tackle the problem.

Now the algorithm is as follows [10]:

1. Each patient points to their preferred GP, if the preferred GP is not in the graph the patient points to its current GP. Each GP points to the patient in their panel with highest priority.
2. Find a cycle in the graph, remove patients part of that cycle. If a GP has no more patients that are in the graph it is removed from the graph.
3. Repeat step 1 until there are no more patients

Like the classical TTC we are guaranteed to have a cycle because of the pigeonhole principle, since if a patient does not get their preferred GP they end up pointing to their own GP making a cycle.

2.3. Cycle cancelling

Cycle cancelling is a technique used in flow and circulation problems to find a minimum cost solution. First we explain some of the terms needed for cycle cancelling then we go onto what cycle cancelling is. We use definitions from Ahuja, Magnanti and Orlin [11].

2.3.1. Minimum cost maximum circulation problem

Let $G = (V, A)$ be a directed multigraph, with n vertices and m arcs. Each arc is a triple (v, w, i) , where v is its tail, w is its head and i is its index. The pair (v, w) specifies the endpoints of the arc, and the index i distinguishes arcs that share the same endpoints, so G may contain parallel arcs.

The multigraph G is a circulation network if each arc (v, w, i) has a capacity $u(v, w, i) \geq 0$ and a cost $c(v, w, i) \in \mathbb{Z}$. For a vertex $w \in V$, let $\text{in}(w) = \{(v, w, i) \in A\}$ denote the set of arcs whose head is w , and $\text{out}(w) = \{(w, v, i) \in A\}$ the set of arcs whose tail is w .

Definition 2.3.1.1 (Circulation): A circulation is a function f assigning a value $f(v, w, i)$ to each arc, subject to the following constraints:

$$0 \leq f(v, w, i) \leq u(v, w, i) \quad \forall (v, w, i) \in A \quad (1)$$

$$\sum_{(u, w, i) \in \text{in}(w)} f(u, w, i) = \sum_{(w, v, j) \in \text{out}(w)} f(w, v, j) \quad \forall w \in A \quad (2)$$

The cost of a circulation f is given by

$$\text{cost}(f) = \sum_{(v, w, i) \in A} c(v, w, i) f(v, w, i) \quad (3)$$

Definition 2.3.1.2 (Minimum cost maximum circulation problem): The minimum cost maximum circulation problem is to take in some circulation network G and find the circulation that has the maximum circulation while minimum cost.

2.3.2.

2.3.3. Algorithm

Using the minimum-cost circulation theorem we can now formalize the cycle cancelling algorithm. Note that we say the starting circulation can be 0 on all edges, sometimes a circulation problem has lower bounds that make this an invalid starting circulation. For cases with lower bounds on edges, a starting circulation can be computed using any max-flow algorithm.

- 1 Start with any feasible circulation f , this can be 0
- 2 **while** $\exists$ negative residual cycle c
- 3 | Cancel c : $\forall e \in c : f(e) := f(e) + \text{capacity}(c)$

Recall that due to the antisymmetry rule of a circulation $f(v, w) = -f(w, v)$ must hold for all edges. When the algorithm cancels a cycle by augmenting flow by $\text{capacity}(c)$ along each edge in the cycle, it implicitly augments $-\text{capacity}(c)$ along the corresponding reverse edges. For example, when flow is pushed by $\text{capacity}(c)$ on an edge (v, w) the flow on its reverse is updated to $f(w, v) = f(w, v) - \text{capacity}(c)$. If we now look at the residual capacity of (w, v) , $u_f(w, v) = u(w, v) - f(w, v)$. Since $f(w, v)$ decreases when $f(v, w)$ increases, pushing flow along an edge increases the residual capacity of its reverse edge. This allows the algorithm to later push flow on the reverse edge to “undo” earlier pushes on certain edges.

2.3.4. Runtime

Finding the negative residual cycles can be done with Bellman-Ford in $O(nm)$ time [12]. At each negative residual cycle that is cancelled the total cost of the circulation decreases by at least 1. The cost is bounded below by $-mCU$ where C is the max cost and U is the max capacity of any edge. It follows that the number of iterations is bounded by $O(mCU)$ and then the runtime is $O(nm^2CU)$. Note that this runtime is pseudo-polynomial since it

depends largely on the size of C and U . Goldberg and Tarjan proved that a variant of the cycle cancelling algorithm called Minimum Mean Cycle Cancelling has a strongly polynomial bound on its runtime [13]. But, we will later show how for our problem we still have a polynomial algorithm using the classical Cycle Cancelling algorithm.

3. Problem Formalization

In this chapter we define the GP allocation problem and also some variations of it. In addition we go through how to construct a graph given a GP allocation problem and finally how we can reduce the GP allocation problem to the Cycle Cover problem in directed graphs.

3.1. The GP allocation problem

3.1.1. Input

In the GP allocation problem we are given a set of patients $P = \{p_1, p_2, \dots, p_n\}$ and a set of doctors $D = \{d_1, d_2, \dots, d_m\}$. We are also given the current and preferred GP assignments, for each patient:

- $D_{\text{cur}}[i]$ denotes the index of the current doctor assigned to patient p_i (or $\perp$ if unassigned)
- $D_{\text{pref}}[i]$ denotes the index of the preferred doctor for patient p_i

In addition we are given a priority function $R : P \rightarrow \mathbb{N}$, mapping some positive integer to each patient, with higher numbers indicating a higher priority to help this patient.

3.1.2. Feasible solution

A feasible solution for the GP allocation problem is a subset of patients $S \subseteq P$ such that the directed graph

$$G_S = (S, E_S), E_S = \{(a, b) \mid a, b \in S, D_{\text{pref}}[\text{idx}(a)] = D_{\text{cur}}[\text{idx}(b)]\} \quad (4)$$

consists of a vertex-disjoint union of directed cycles that cover all vertices in S . This means that S is a set of patients that can all get their preferred doctor if we allow for simultaneous exchanges following cycles.

3.1.3. Optimization criterion

Next we can start defining variants of feasible solutions where we want to maximize some *score*. Examples of such a score could be to find a feasible solution with many patients or something based on priority. We might want a solution that exchanges as many patients as possible, exchanges patients with the highest priority or some mix of these.

First we define our general ordering

Definition 3.1.3.1 (Optimization criterion): An ordering $\succ$ over feasible solutions. A solution is optimal if it is maximal under $\succ$.

Build on this to create our three main variants of scoring functions.

3.1.3.1. Lexicographic maximization by priority

Definition 3.1.3.1.1 (Characteristic vector): Let patients be ordered by priority: $p_1 \geq p_2 \geq \dots \geq p_n$ where p_1 has the highest priority. Given a feasible solution $S \subseteq P$, the **characteristic vector** is:

$$\chi(S) = (b_1, b_2, \dots, b_n) \in \{0, 1\}^n, \quad b_i = \begin{cases} 1 & \text{if } p_i \in S \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

Definition 3.1.3.1.2 (Lexicographic ordering by priority):

$$S \succ_{\text{lex}} S' \quad \text{iff} \quad \chi(S) \text{ is lexicographically greater than } \chi(S') \quad (6)$$

That is, at the first index i where S and S' differ, $\chi(S)_i = 1$ and $\chi(S')_i = 0$.

Intuitively, $\chi(S)$ is a binary number whose most significant non-zero bit corresponds to the highest-priority patient that gets helped. The optimal solution is then the one that maximises this binary number: always satisfying the highest-priority patient it can, then subject to that the next highest, and so on.

Example (Ordering two solutions lexicographically by priority):

We use the example graph in Figure 4 as the graph to create solutions from.

Given solutions

1. $S = \{p_4, p_5\}$ represents the subgraph G_S containing the cycle $p_4 \rightarrow p_5 \rightarrow p_4$
2. $S' = \{p_1, p_2, p_3\}$ represents the subgraph $G_{S'}$ containing the cycle $p_1 \rightarrow p_2 \rightarrow p_3 \rightarrow p_1$

For simplicity's sake we say that $R(p_i) = i$, so p_5 has the highest priority.

Ordering all five patients by priority gives $p^{(1)} = p_5, p^{(2)} = p_4, \dots, p^{(5)} = p_1$. Then:

$$\chi(S) = (1, 1, 0, 0, 0) \quad \chi(S') = (0, 0, 1, 1, 1) \quad (7)$$

At the first differing position ($i = 1$), $\chi(S)_1 = 1 > 0 = \chi(S')_1$, so $S \succ_{\text{lex}} S'$.

So for our first variant we want to find the maximal solution under the $\succ_{\text{lex}}$ ordering. This means always prioritizing patients with highest priority.

3.1.3.2. Maximizing cardinality

Another natural variant is to find a solution with the largest cardinality, e.g. a solution that contains the most patients.

Definition 3.1.3.2.1 (Ordering by cardinality):

$$S \succ_{\text{size}} S' \quad \text{iff} \quad |S| > |S'| \quad (8)$$

This optimal solution will then be one that exchanges the most patients.

Example (Ordering two solutions by cardinality):

Using the same solutions as in the previous example:

1. $S = \{p_4, p_5\}$ represents the subgraph G_S containing the cycle $p_4 \rightarrow p_5 \rightarrow p_4$
2. $S' = \{p_1, p_2, p_3\}$ represents the subgraph $G_{\{S'\}}$ containing the cycle $p_1 \rightarrow p_2 \rightarrow p_3 \rightarrow p_1$

Under $\succ_{\text{size}} : |S| = 2$ and $|S'| = 3$, so $S' \succ_{\text{size}} S$. However the solution is $\{p_1, p_2, p_3, p_4, p_5\}$, the union of both cycles.

Notice that the same two sets are reversibly ordered under the two criteria: $S \succ_{\text{lex}} S'$ (because S contains the highest-priority patient p_5) but $S' \succ_{\text{size}} S$ (because S' has more patients).




Figure 4: An example graph G .

3.1.3.3. Maximizing utility

Finally we formulate an ordering that is a combination of $\succ_{\text{lex}}$ and $\succ_{\text{size}}$. We define the total utility to be the sum of the priorities of the patients in S .

We recall that $R : P \rightarrow \mathbb{N}$ is the function mapping a patient to a priority.

Definition 3.1.3.3.1 (Total utility function):

$$U(S) = \sum_{p \in S} R(p) \tag{9}$$

Definition 3.1.3.3.2 (Ordering by total utility):

$$S \succ_{\text{util}} S' \text{ iff } U(S) > U(S') \quad (10)$$

This means that even though solution S' contains some high priority patients, if S contains enough lower priority patients then S can still be a higher ranked solution under $\succ_{\text{util}}$.

Example (Ordering two solutions by total utility):

We use the following solutions from Figure 5:

1. $S = \{1, 3, 4\}$ representing the cycle $p_1 \rightarrow p_3 \rightarrow p_4 \rightarrow p_1$
2. $S' = \{5, 2\}$ representing the cycle $p_5 \rightarrow p_2 \rightarrow p_5$

This gives the following total utilities:

$$U(S) = 1 + 3 + 4 = 8$$

$$U(S') = 2 + 5 = 7$$

It follows that $S \succ_{\text{util}} S'$. Here we see that the solution with more patients with lower maximum priority has a higher score than another with greater maximum priority.




Figure 5: An example graph G.

3.1.4. Priority function

The priority function has quite a large effect on the final solution. The priority chosen will largely decide whether that patient will be in the final solution or not. This is why its important to discuss what effects the function can have on solutions when we are maximizing the ordering $\succ_{\text{util}}$ or $\succ_{\text{lex}}$.

If we make the priority function exponential such that $R(a) = 2^a$ then using the ordering $\succ_{\text{util}}$ becomes equal to $\succ_{\text{lex}}$. This is because if we have for patient a priority of 2^a then even if we sum all While this ordering is in terms maybe the most fair, since we never

let a patient with lesser priority switch before one with greater, it might not always be the best for the collective good, since a high priority patient might block a lot of lower priority patients from switching. This way we might want a priority function like $R(a) = \text{Days patient } a \text{ has been waiting for a switch}$. This way if we have the choice between patient p_i who has waited for 30 days $R(p_i) = 30$ and 4 patients who have waited 40 days in total an algorithm maximizing $\succ_{\text{util}}$ will choose the 4 patients.

We can adjust the priority function to prioritize higher priority by making it exponential but still under the 2^a . When we use days waited we can make $R(a) = 1.1^{\text{Days patient } a \text{ has been waiting}}$. Adjusting this closer to 2 makes higher priority patients more prioritized and vice versa for lowering it.

4. Implementation

In this chapter we look at each algorithm we have implemented, how we implemented them, why and runtime analysis. We have in total implemented 4 algorithms, Greedy DFS, Cycle Cancellling for Cardinality $\succ_{\text{size}}$, Cycle Cancellling for Utility $\succ_{\text{util}}$ and Cycle Cancellling for Lexicographic Priority $\succ_{\text{lex}}$. For the cycle cancelling algorithms we also give, or refer to existing, proofs for why they are exact.

4.1. Different graph representations used

In the algorithms we use different graph representations, so we start by defining the different graph representations of *The GP allocation problem*.

4.1.1. Patient and Doctor graph

First we have the easiest graph containing both patients and doctors as vertices. The edges are directed from a patient to a doctor and from a doctor to a patient. Edges never go patient $\rightarrow$ patient or doctor $\rightarrow$ doctor. An edge patient $\rightarrow$ doctor symbolizes that the patient has that doctor as its preferred doctor, and doctor $\rightarrow$ patient symbolises that the doctor currently has that patient as one of his or hers current patients. Observe that this graph is bipartite as we have patients on one side and doctors on the other. Now the formal definition of our graph. Let $I = \{0, \dots, |P| - 1\}$. Then:

$$\begin{aligned} G &= (V, E), \quad V = P \cup D \\ E &= \left\{ (p_i, D_{\text{pref}[i]}) \mid i \in I \right\} \cup \left\{ (D_{\text{cur}[i]}, p_i) \mid i \in I \right\} \end{aligned} \quad (11)$$

4.1.2. Doctor graph collapsed edges

In this weighted graph we condense the problem to only have doctors as nodes and edges between doctors. An edge doctor $a \rightarrow$ doctor b symbolises that there exists a patient that wants to switch from doctor a to doctor b , or that the patient currently has doctor a and has doctor b as preferred. The capacity of an edge indicates the number of patients wanting that switch, while the cost is -1 . We define our graph as:

$$\begin{aligned} G &= (V, E), \quad V = D \\ E &= \left\{ (a, b), (b, a) \mid \exists i \in I \text{ s.t. } D_{\text{cur}[i]} = a \wedge D_{\text{pref}[i]} = b \right\} \\ u(a, b) &= \left| \left\{ i \in I \mid D_{\text{cur}[i]} = a \wedge D_{\text{pref}[i]} = b \right\} \right| \\ u(b, a) &= 0 \\ c(a, b) &= -1 \\ c(b, a) &= 1 \end{aligned} \quad (12)$$

4.1.3. Doctor graph priority weighted

This weighted graph is much like the Doctor graph collapsed edges, but while that graph representation focuses on number of patients wanting a switch this graph focuses on the priority of each patient. Instead of collapsing all preferences that are equal here we make each preference into an edge and weight it with the priority of that patient.

$$\begin{aligned}
G &= (V, E), \quad V = D \\
E &= \left\{ \left(D_{\text{cur}[i]}, D_{\text{pref}[i]}, i \right) \mid i \in I \right\} \\
u(a, b, i) &= 1 \\
c(a, b, i) &= -R(P_i)
\end{aligned} \tag{13}$$

4.2. Greedy DFS

We first consider a greedy approach inspired by the Top Trading Cycles implementation of Huitfeldt et al. Their algorithm preserves TTC properties at each round by restricting each doctor node to a single outgoing edge. We relaxed this constraint, allowing each doctor to maintain outgoing edges to all of its current patients simultaneously, and resolved ties by patient priority.

We are given the lists of patients and doctors P, D , the assignment arrays $D_{\text{cur}}, D_{\text{pref}}$, and a priority function $R : P \rightarrow \mathbb{N}$, where a higher number means higher priority. We model the problem as a bipartite directed graph over patients and doctors. Let $I = \{0, \dots, |P| - 1\}$. Then:

$$G = (V, E), \quad V = P \cup D, \quad E = \left\{ \left(p_i, D_{\text{pref}[i]} \right) \mid i \in I \right\} \cup \left\{ \left(D_{\text{cur}[i]}, p_i \right) \mid i \in I \right\}$$

Each patient p_i has an outgoing edge to their preferred doctor, and each doctor has outgoing edges to all of its currently registered patients. A directed cycle in G necessarily alternates between patient and doctor nodes and corresponds to a valid exchange: each patient in the cycle moves to their preferred doctor, and each doctor loses exactly one patient while gaining one.

The algorithm processes patients in decreasing priority order. For each patient, a DFS attempts to find a cycle through that patient; when the DFS reaches a doctor node with multiple candidate patients, it always explores the highest-priority one first.

```

1 let resolved_patients = []
2 let p_prio = Sorted list of patients by priority
3 let wants_to_switch = [true] * |P|
4 for each p in p_prio do
5   if let cycle = dfs_find_cycle(G, R, p) then
6     for each p in cycle do
7       let i = idx(p)
8       wants_to_switch[i] = false
9       resolved_patients.push(p)
10    end
11  end
12 end
13 return resolved_patients

```

The greedy rule ensures that high-priority patients are preferentially included in cycles, but it does not guarantee a globally optimal solution. The following example illustrates how the greedy choice can be locally motivated yet globally suboptimal.




Figure 6: Example graph G where Greedy DFS would choose suboptimal solution. px means patient x and has priority x .

In Figure 6 each patient px has priority x . The DFS begins at $p4$ (highest priority), follows the edge to $d2$, and there chooses $p2$ over $p1$ since $R(p2) > R(p1)$. The resulting cycle $p4 \rightarrow d2 \rightarrow p2 \rightarrow d1 \rightarrow p4$ is committed, leaving $p1$ and $p3$ unsatisfied with no further cycles remaining.

Had the DFS chosen $p1$ at $d2$ instead, it would have found the longer cycle $p4 \rightarrow d2 \rightarrow p1 \rightarrow d3 \rightarrow p3 \rightarrow d1 \rightarrow p4$, satisfying three patients. The greedy choice at $d2$ was locally motivated by priority but globally suboptimal. This motivates the exact algorithms in the following sections.

5. Experimental Results

5.1. Experimental Setup

5.2. Performance Comparison

5.3. Impact of Priority Strategies

6. Conclusion

6.1. Summary

6.2. Future Work

Bibliography

- [1] Legelisten, “Med rett til å bytte – Ventetider for å bytte fastlege i Norge.” Accessed: Apr. 08, 2026. [Online]. Available: <https://legelisten.s3.eu-west-1.amazonaws.com/blog-files/Med+rett+til+%C3%A5+bytte++Ventetider+for+%C3%A5+bytte+fastlege+i+Norge.pdf>[◦]
- [2] Helfo, “Fastlegestatistikk.” Accessed: Apr. 08, 2026. [Online]. Available: <https://www.helfo.no/Fastlegeordninga/fastlegestatistikk>[◦]
- [3] Statistisk sentralbyrå, “Fastlegelister og fastlegekonsultasjoner, etter år og region (12005).” Accessed: Apr. 08, 2026. [Online]. Available: <https://www.ssb.no/statbank/table/12005>[◦]
- [4] L. Shapley and H. Scarf, “On cores and indivisibility,” *Journal of Mathematical Economics*, vol. 1, no. 1, pp. 23–37, 1974, doi: [https://doi.org/10.1016/0304-4068\(74\)90033-0](https://doi.org/10.1016/0304-4068(74)90033-0)[◦].
- [5] D. J. Abraham, A. Blum, and T. Sandholm, “Clearing algorithms for barter exchange markets: enabling nationwide kidney exchanges,” pp. 295–304, 2007, doi: [10.1145/1250910.1250954](https://doi.org/10.1145/1250910.1250954)[◦].
- [6] Wikipedia contributors, “Optimal kidney exchange — Wikipedia, The Free Encyclopedia.” Accessed: May 05, 2026. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Optimal_kidney_exchange&oldid=1351916222[◦]
- [7] D. Gale and L. S. Shapley, “College Admissions and the Stability of Marriage,” *The American Mathematical Monthly*, vol. 69, no. 1, pp. 9–15, 1962, Accessed: May 06, 2026. [Online]. Available: <http://www.jstor.org/stable/2312726>[◦]
- [8] Wikipedia contributors, “Top trading cycle — Wikipedia, The Free Encyclopedia.” Accessed: May 08, 2026. [Online]. Available: https://en.wikipedia.org/w/index.php?title=Top_trading_cycle&oldid=1344910705[◦]
- [9] A. E. Roth, “Incentive compatibility in a market with indivisible goods,” *Economics Letters*, vol. 9, no. 2, pp. 127–132, 1982, doi: [https://doi.org/10.1016/0165-1765\(82\)90003-9](https://doi.org/10.1016/0165-1765(82)90003-9)[◦].
- [10] I. Huitfeldt, V. Marone, and D. C. Waldinger, “Designing Dynamic Reassignment Mechanisms: Evidence from GP Allocation,” Working Paper 32458, May 2024. doi: [10.3386/w32458](https://doi.org/10.3386/w32458)[◦].
- [11] R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, *Network Flows: Theory, Algorithms, and Applications*. Prentice Hall, 1993.
- [12] Wikipedia contributors, “Bellman–Ford algorithm — Wikipedia, The Free Encyclopedia.” [Online]. Available: https://en.wikipedia.org/w/index.php?title=Bellman%E2%80%93Ford_algorithm&oldid=1351231946[◦]
- [13] A. V. Goldberg and R. E. Tarjan, “Finding minimum-cost circulations by canceling negative cycles,” *J. ACM*, vol. 36, no. 4, pp. 873–886, Oct. 1989, doi: [10.1145/76359.76368](https://doi.org/10.1145/76359.76368)[◦].

Index of Figures

Figure 1	An example graph for the kidney exchange problem.	8
Figure 2	An example Housing Market for TTC.	9
Figure 3	An example Housing Market for TTC, A's $\text{Top}(i)$ is itself.	9
Figure 4	An example graph G.	15
Figure 5	An example graph G.	16
Figure 6	Example graph G where Greedy DFS would choose suboptimal solution. px means patient x and has priority x	20